

Lecture 12

Support Vector Machines

ISLR 9



Jilles Vreeken
Aleksandar Bojchevski



UNIVERSITÄT
DES
SAARLANDES




CISPA
HELMHOLTZ CENTER FOR
INFORMATION SECURITY

Hyperplanes

- in p -dimensional vector space, a linear hyperplane is a $(p-1)$ -dimensional subspace
- equivalently, a linear hyperplane is the set of points that satisfy a linear equation of the form $\beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$
- an **affine** hyperplane is the set of points that fulfills $\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$ for some $\beta_0 \neq 0$
- a hyperplane divides the vector space into two half spaces
- the vector $(\beta_1, \dots, \beta_p)$ is the **normal** vector of the hyperplane

the hyperplane $f(X) = 1 + 2X_1 + 3X_2 = 0$




Classification Using Separating Hyperplanes

- assume a data matrix $\mathbf{x}_1 = \begin{pmatrix} x_{11} \\ \vdots \\ x_{1p} \end{pmatrix}, \dots, \mathbf{x}_n = \begin{pmatrix} x_{n1} \\ \vdots \\ x_{np} \end{pmatrix}$
for a binary classification problem with classes $\{1, -1\}$
- assume further a test vector $\mathbf{x}^* = (x_1^*, \dots, x_p^*)^T$

We define a classifier based on a **separating hyperplane**

- the data points of the two classes locate in separate half spaces

different separating hyperplanes




Classification Using Separating Hyperplanes

- assume a data matrix $\mathbf{x}_1 = \begin{pmatrix} x_{11} \\ \vdots \\ x_{1p} \end{pmatrix}, \dots, \mathbf{x}_n = \begin{pmatrix} x_{n1} \\ \vdots \\ x_{np} \end{pmatrix}$
for a binary classification problem with classes $\{1, -1\}$
- assume further a test vector $\mathbf{x}^* = (x_1^*, \dots, x_p^*)^T$

We define a classifier based on a **separating hyperplane**

- the data points of the two classes locate in separate half spaces
- the hyperplane is defined by $\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$
- the classification is $\text{sign}(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p)$

separating hyperplane and resulting classifier




- the distance of a point from the hyperplane is informative about the confidence in the classification

The Maximal Margin Classifier

- a hyperplane that maximizes the **distance of the closest point** in the training set to it can be considered optimal

separating hyperplane and resulting classifier




The Maximal Margin Classifier

- a hyperplane that maximizes the **distance of the closest point** in the training set to it can be considered optimal
- this distance is called the **margin**

The closest data points are called the **support vectors**

- only they determine the hyperplane
- can be a small subset of all points

separating hyperplane and resulting classifier




Constructing the Maximal Margin Classifier

The optimization problem is

$$\begin{aligned} & \max_{\beta_0, \beta_1, \dots, \beta_p, M} M \\ & \text{subject to } \sum_{j=1}^p \beta_j^2 = 1 \quad \text{normal vector is unit vector} \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M, \quad i = 1, \dots, n \quad \text{correct classification, if } M > 0 \end{aligned}$$

- since the normal is a unit vector, the distance of point i from the hyperplane is given by $y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip})$
- solve the optimization problem with convex optimiz. techniques
- often, there is no separating hyperplane

separating hyperplane and resulting classifier




Constructing the Maximal Margin Classifier

The optimization problem is

$$\begin{aligned} & \max_{\beta_0, \beta_1, \dots, \beta_p, M} M \\ & \text{subject to } \sum_{j=1}^p \beta_j^2 = 1 \quad \text{normal vector is unit vector} \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M, \quad i = 1, \dots, n \quad \text{correct classification, if } M > 0 \end{aligned}$$

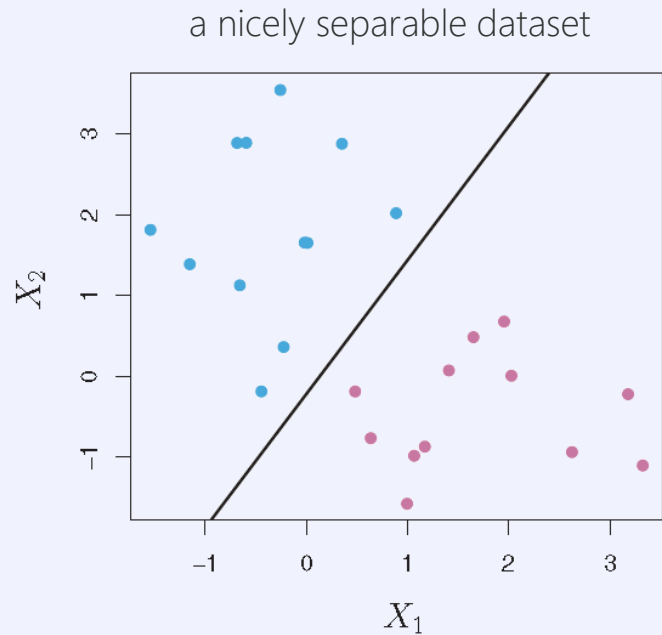
- since the normal is a unit vector, the distance of point i from the hyperplane is given by $y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip})$
- solve the optimization problem with convex optimiz. techniques
- often, there is no separating hyperplane
- then we have to generalize to allow for misclassifications

a non-separable dataset



The Support Vector Classifier

Even if the dataset is separable, a separating hyperplane may not be desirable



The Support Vector Classifier


Even if the dataset is separable, a separating hyperplane may not be desirable

- adding a single data point leads to a hard-to-separate dataset
- the classifier is extremely sensitive to changes in the data

Sometimes it may be preferable to have a classifier that misplaces a few points in the training set but has a large margin to the other data points

- greater robustness w.r.t to small changes in the data
- better classification of most of the training points
- the **soft-margin classifier** does exactly this

a nicely separable dataset with an outlier




The Support Vector Classifier

Even if the dataset is separable, a separating hyperplane may not be desirable

- adding a single data point leads to a hard-to-separate dataset
- the classifier is extremely sensitive to changes in the data

Sometimes it may be preferable to have a classifier that misplaces a few points in the training set but has a large margin to the other data points

- greater robustness w.r.t to small changes in the data
- better classification of most of the training points
- the **soft-margin classifier** does exactly this




The Support Vector Classifier

Even if the dataset is separable, a separating hyperplane may not be desirable

- adding a single data point leads to a hard-to-separate dataset
- the classifier is extremely sensitive to changes in the data

Sometimes it may be preferable to have a classifier that misplaces a few points in the training set but has a large margin to the other data points

- greater robustness w.r.t to small changes in the data
- better classification of most of the training points
- the **soft-margin classifier** does exactly this



The Support Vector Classifier


Even if the dataset is separable, a separating hyperplane may not be desirable

- adding a single data point leads to a hard-to-separate dataset
- the classifier is extremely sensitive to changes in the data

Sometimes it may be preferable to have a classifier that misplaces a few points in the training set but has a large margin to the other data points

- greater robustness w.r.t to small changes in the data
- better classification of most of the training points
- the **soft-margin classifier** does exactly this

- points can be on the wrong side of the margin (misplaced but correct) or the hyperplane (misclassified)



Details of Soft-Margin Support Vector Classifier


The optimization problem is now

$$\begin{aligned} & \max_{\beta_0, \beta_1, \dots, \beta_p, M} M \\ & \text{subject to } \sum_{j=1}^p \beta_j^2 = 1 \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\ & \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \end{aligned}$$

Slack variables allow for a fractional violation of the hard margin constraint
↓
Budget for total admissible misclassification

The following holds if we also choose the smallest possible ϵ_i :

- $\epsilon_i = 0 \Rightarrow$ the observation is on the **correct** side of the **margin**
- $\epsilon_i > 0 \Rightarrow$ the observation is on the **wrong** side of the **margin**
- $\epsilon_i > 1 \Rightarrow$ the observation is on the **wrong** side of the **hyperplane**
- Furthermore: no more than C observations can be on the wrong side of the hyperplane

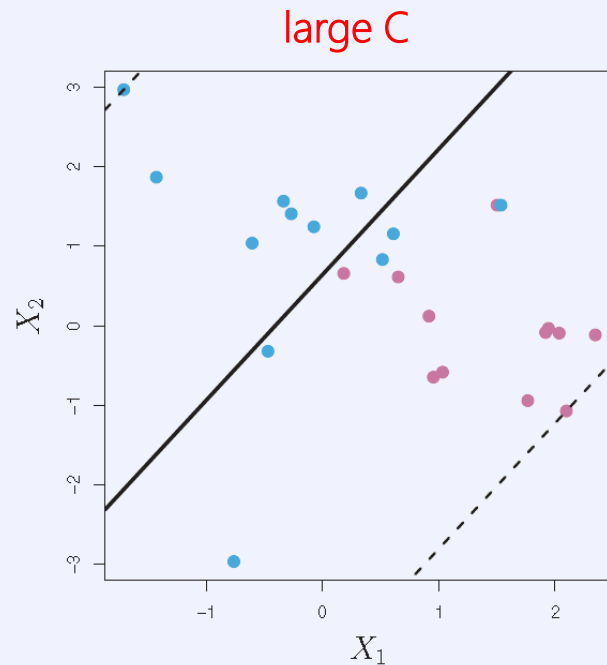


On the Effect of the Budget \mathcal{C}

The optimization problem is now

$$\begin{aligned} & \max_{\beta_0, \beta_1, \dots, \beta_p, M} M \\ & \text{subject to } \sum_{j=1}^p \beta_j^2 = 1 \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\ & \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq \mathcal{C} \end{aligned}$$

As \mathcal{C} increases, we become more tolerant to violations



On the Effect of the Budget \mathcal{C}

The optimization problem is now


$$\max_{\beta_0, \beta_1, \dots, \beta_p, M}$$

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq \mathcal{C}$$

As \mathcal{C} increases, we become more tolerant to violations




On the Effect of the Budget \mathcal{C}

The optimization problem is now

$$\begin{aligned} & \max_{\beta_0, \beta_1, \dots, \beta_p, M} M \\ & \text{subject to } \sum_{j=1}^p \beta_j^2 = 1 \\ & y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\ & \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq \mathcal{C} \end{aligned}$$

As \mathcal{C} increases, we become more tolerant to violations



On the Effect of the Budget \mathcal{C}

The optimization problem is now

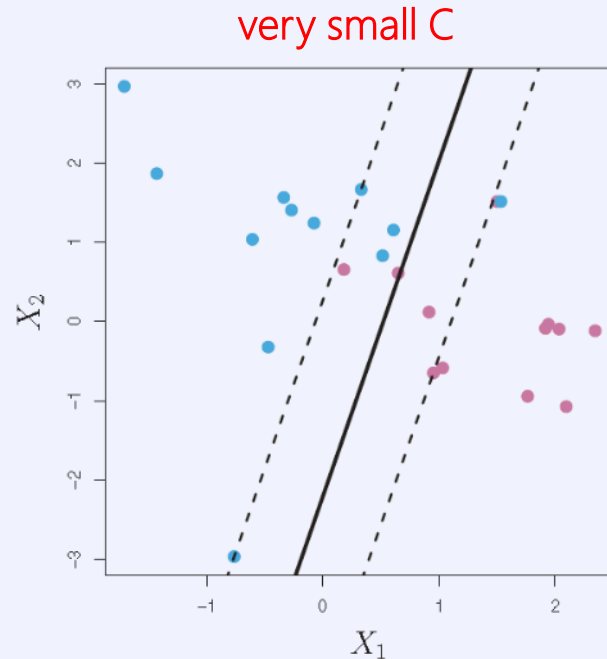
$$\max_{\beta_0, \beta_1, \dots, \beta_p, M} M$$

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq \mathcal{C}$$

As \mathcal{C} increases, we become more tolerant to violations



The Margin and the Support Vectors


We choose C via cross-validation

For the soft-margin classifier support vectors all either lie exactly on the margin or on the wrong side of the margin

- intuition: since only changing those points would affect the hyperplane
- C controls the bias-variance tradeoff
- with large C the margin is wide and there are many support vectors
 - low variance and potentially high bias
- with small C the margin is thin and there are a few support vectors
 - high variance and small bias

The fact that correctly classified points far away from the hyperplane do not affect the classifier is a property of the support-vector classifier

as C decreases we become less tolerant to violations



Nonlinear Decision Boundaries

Sometimes, data is inherently nonlinear

- then there is no soft margin that will do the trick
- we need a nonlinear version of support vector machines
- we could add nonlinear features to the feature space, e.g. $X_1, X_1^2, X_2, X_2^2, \dots, X_p, X_p^2$ instead of X_1, X_2, \dots, X_p
- the resulting optimization program would become


$$\max_{\beta_0, \beta_{11}, \beta_{12}, \dots, \beta_{p1}, \beta_{p2}, \epsilon_1, \dots, \epsilon_n} M$$

$$\text{subject to } \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C, \sum_{j=1}^p \sum_{k=1}^2 \beta_{jk}^2 = 1$$

$$y_i \left(\beta_0 + \sum_{j=1}^p \beta_{j1} x_{ij} + \sum_{j=1}^p \beta_{j2} x_{ij}^2 \right) \geq M(1 - \epsilon_i), \quad i = 1, \dots, n$$

- we could add higher-order, interaction terms, or use functions other than polynomials

the true boundary is non-linear



The Kernel Trick

With **support vectors machines (SVMs)** there is a different very elegant trick – the **kernel trick**

- builds on the optimization procedure for SVMs, which we will not detail
- it suffices to say that the linear support vector classifier can be rewritten as $f(x^*) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x^*, x_i \rangle$
- $\langle x^*, x_i \rangle = \sum_{j=1}^p x_j^* x_{ij}$ is the inner product,
- and the α_i are parameters that result from the training

set of support vectors

Important: Only the α_i for the support vectors are nonzero $f(x^*) = \beta_0 + \sum_{i \in S} \alpha_i \langle x^*, x_i \rangle$



The Kernel Trick

Only the α_i for the support vectors are nonzero $f(x^*) = \beta_0 + \sum_{i \in S} \alpha_i \langle x^*, x_i \rangle$

- to calculate α_i and β_0 we only need $\frac{n(n-1)}{2}$ inner products $\langle x_i, x_{i'} \rangle$ between all pairs of training points
- the actual coordinates of the training observations or the test point are never needed!

We can generalize inner products to (nonlinear) kernels $K(x_i, x_{i'})$

- a kernels quantifies the similarity between two data points
- a simple linear kernel is $K(x_i, x_{i'}) = \langle x_i, x_{i'} \rangle$
- it quantifies similarity in terms of the standard (Pearson) correlation

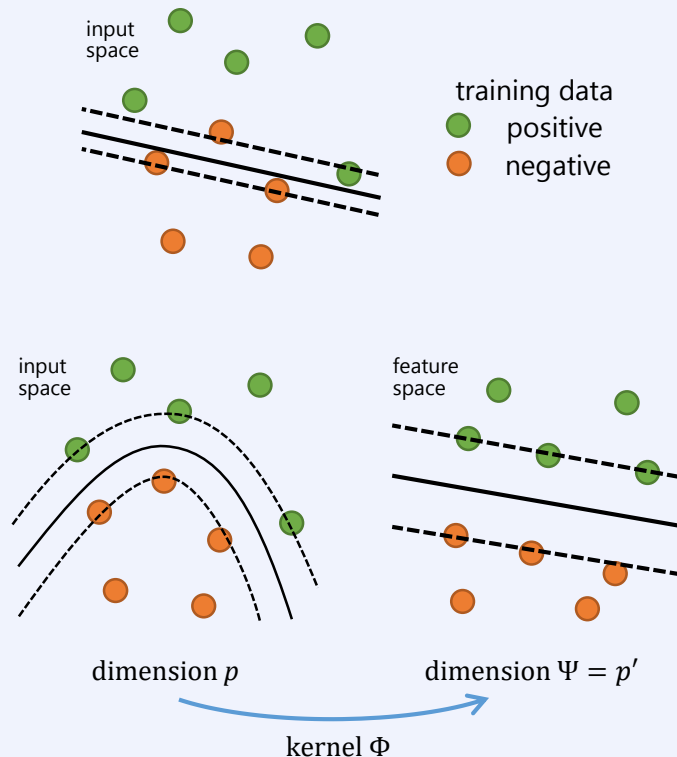


Nonlinear Kernels

Two popular choices:

- The polynomial kernel with degree d
$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p \langle x_{ij}, x_{i'j} \rangle\right)^d$$
- The radial-basis kernel
$$K(x_i, x_{i'}) = \exp\left(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2\right)$$
- in general, a kernel is any symmetric and positive definite function¹ of its two arguments

A **VERY** important theorem says that for any kernel K there is a function $\Phi: \mathbb{R}^p \rightarrow \Psi$ such that $K(x_i, x_j) = \langle \Phi(x_i), \Phi(x_j) \rangle$

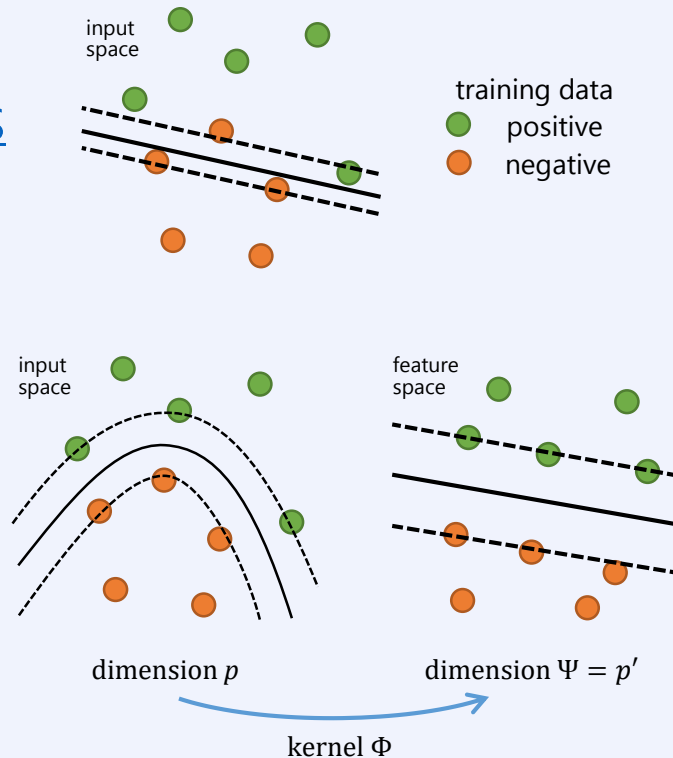





Reproducing Kernel Hilbert Space (RKHS)

Applying the kernel actually means performing an inner product in some space Ψ , the so-called [RKHS](#)

- neither Φ nor Ψ generally can (or need) be constructed in a computationally usable form
- however in some cases, they can, e.g., for $p = 2$ and the polynomial kernel with $d = 2$, we have $\dim \Psi = p' = 6$ and
$$\begin{aligned}\Phi_1(X) &= 1, & \Phi_2(X) &= \sqrt{2}X_1, & \Phi_3(X) &= \sqrt{2}X_2 \\ \Phi_4(X) &= X_1^2, & \Phi_5(X) &= X_2^2, & \Phi_6(X) &= \sqrt{2}X_1X_2\end{aligned}$$
- for the radial basis kernel, p' is infinite¹




The Radial Basis Kernel

If our test point \mathbf{x}^* is far from the training point \mathbf{x}_i then $\sum_{j=1}^p (\mathbf{x}_j^* - \mathbf{x}_{ij})^2$ will be large, so the kernel value $\exp\left(-\gamma \sum_{j=1}^p (\mathbf{x}_j^* - \mathbf{x}_{ij})^2\right)$ will be tiny

- thus \mathbf{x}_i will not influence the value of $f(\mathbf{x}^*)$ by much

Since the class label is based on the sign of $f(\mathbf{x}^*)$ the radial basis kernel thus has very local behavior

- γ controls the locality
- decreasing γ increases locality




The Radial Basis Kernel

If our test point \mathbf{x}^* is far from the training point \mathbf{x}_i then $\sum_{j=1}^p (\mathbf{x}_j^* - \mathbf{x}_{ij})^2$ will be large, so the kernel value $\exp\left(-\gamma \sum_{j=1}^p (\mathbf{x}_j^* - \mathbf{x}_{ij})^2\right)$ will be tiny

- thus \mathbf{x}_i will not influence the value of $f(\mathbf{x}^*)$ by much

Since the class label is based on the sign of $f(\mathbf{x}^*)$ the radial basis kernel thus has very local behavior

- γ controls the locality
- decreasing γ increases locality



Advantages of Kernels

To calculate the SVM you only need the kernel matrix for the pairs of training points

- in contrast, enlarging the feature space is computationally expensive

Can be applied to arbitrary observations that are not vectors: graphs, strings, molecules, etc.

The kernel trick can also be used with other statistical learning methods such as LDA or PCA

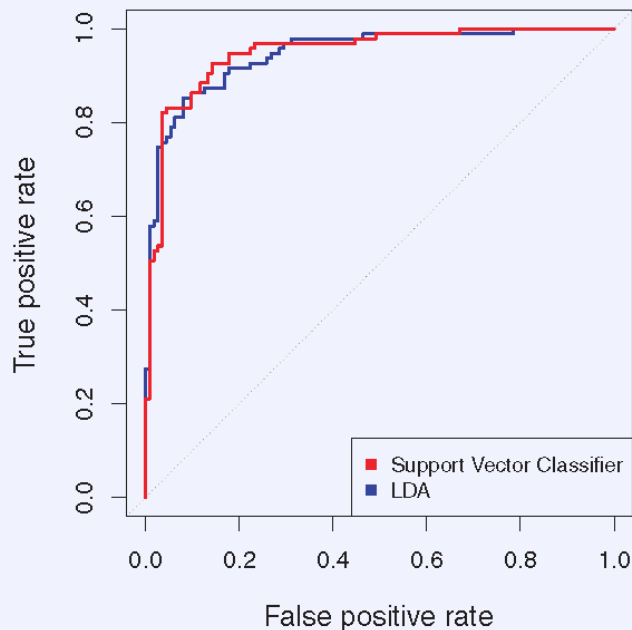
Application to the Heart Disease Data

- 13 predictors are used for classification
- binary target: whether an individual has heart disease
- 207 training, 90 test observations

Comparison of LDA and linear SVM

- use a threshold on $f(\mathbf{x})$ to parameterize SVM
- use a threshold on the linear discriminant to parameterize LDA
- similar performance on the training data

ROC curve for classification performance on the **Heart** dataset – training data




Application to the Heart Disease Data

- 13 predictors are used for classification
- binary target: whether an individual has heart disease
- 207 training, 90 test observations

Comparison of LDA and linear SVM

- use a threshold on $f(\mathbf{x})$ to parameterize SVM
- use a threshold on the linear discriminant to parameterize LDA
- similar performance on the training data
- SVM outperforms LDA on the test set – generalizes better

ROC curve for classification performance on the **Heart** dataset – test data




Application to the Heart Disease Data

- 13 predictors are used for classification
- binary target: whether an individual has heart disease
- 207 training, 90 test observations

Comparison of linear and nonlinear (radial basis kernel) support vector classifiers

- $\gamma = 10^{-1}$ is best on the training set

ROC curve for classification performance on the **Heart** dataset – training data




Application to the Heart Disease Data

- 13 predictors are used for classification
- binary target: whether an individual has heart disease
- 207 training, 90 test observations

Comparison of linear and nonlinear (radial basis kernel) support vector classifiers

- $\gamma = 10^{-1}$ is best on the training set
- $\gamma = 10^{-1}$ is worst on the training set
- this amounts to a very local kernel which incurs high variance
- other nonlinear kernels perform comparably with the linear kernel

ROC curve for classification performance on the **Heart** dataset – test data



Relationship to Logistic Regression

The SVM optimization problem can be rewritten as

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \left\{ \sum_{i=1}^n \max[0, 1 - y_i f(x_i)]_+ + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$


- this has a general form of a regularized regression

$$\min_{\beta_0, \beta_1, \dots, \beta_p} \{L(\mathbf{X}, \mathbf{y}, \beta) + \lambda P(\beta)\} \text{ with loss } L \text{ and penalty } P$$

SVM uses the same penalty as in ridge regression, but a different loss function, called **hinge** loss

- similar to that used in logistic regression, thus both classifiers often give similar results
- with better separation, SVM is better, with more overlap logistic regression tends to be better

The budget C for margin violations is inversely proportional to the penalty parameter λ






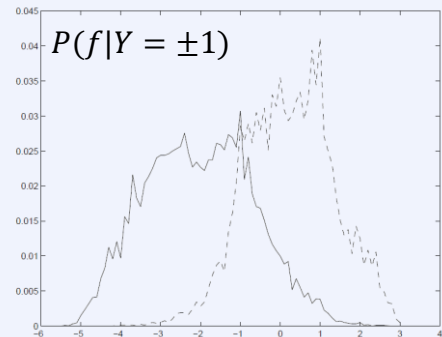
Posterior Probabilities from SVMs

Turning SVM output into ROC curves

- compute posterior probability of the input belonging to class 0 and 1 respectively using the formula
$$P(y = 1|x) = \frac{1}{1 + \exp(Af(x) + B)}$$
 where $f(x)$ is the SVM output
- A and B are parameters that are trained discriminatively

The original distributions are not Gaussian and ragged

- but, logistic fit works well





Multiclass classification

Standard SVM cannot handle multiple classes. We show strategies to address the issue.

- they can be generally applied anytime a binary classifier is the only option

One-vs-rest: Train K SVM models for K classes, where each SVM is being trained for classification of one class against all the remaining ones.

- winner is then the class, where the distance from the hyperplane is maximal

One-vs-one: train $\binom{K}{2}$ classifiers (all possible pairings) and evaluate all

- winner is the class with the majority vote
- votes can be weighted according to the distance from the margin


One-class SVM: an unsupervised algorithm to learn a decision function for novelty detection



Support Vector Machines for Regression

Want to fit a linear model $f(x) = x^T \beta + \beta_0$ such that all data points lie inside a margin of width ϵ of the regression hyperplane

- impose a square penalty on model complexity



The loss function is the ϵ -insensitive $V_\epsilon(r)$

- only data points **on or outside** the tube change the model (this is different from classification)
- these are the support vectors
- kernels distort the tube




Example SVR with a Radial Basis Kernel

SVR with RBF kernel on synthetic data

- Green lines show the ϵ -boundaries
- Blue points represent data instance
- Marked blue points are the support vectors

The fitted model adapts well to the structure of the data

Introducing new datapoints change the model only if they are on the ϵ -boundary or outside of it



Summary

The main ideas behind SVMs is to find the max-margin hyperplane that separate the data

Hard SVM requires that all training data is correctly separated by can overfit

Soft SVM allows violations of the margin up to a budget C to get a better hyperplane overall

We can rewrite the SVM classifier only in terms of inner products – replacing those with a kernel is the kernel trick which allow us to efficiently introduce non-linearity

- the kernel trick is an important general idea that also applies to LDA, PCA and other models

Linear SVM is similar to logistic ridge regression but uses a hinge loss instead